

Part 1: S&P 500 Stock Price Forecast

Motivation: Beginning in March 2020 during the COVID pandemic, the stock market caught an unexpected upwind. I want to know if it would have been possible to predict this trend early on.

```
In [123]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import plotly.graph_objects as go
import plotly.express as px
import datetime
import yfinance as yf
from scipy import stats

from sklearn.preprocessing import MinMaxScaler
from sklearn.linear_model import LinearRegression
import keras.backend as K
from keras.models import Sequential
from keras.layers import LSTM
from keras.layers import Dropout
from keras.layers import Dense
from keras.callbacks import EarlyStopping
```

I) EDA

```
In [124]: def candlestick(df, title, begin_year='1975'):
# Plot candlestick chart
temp = df[df.index > begin_year]
fig = go.Figure(data=[go.Candlestick(x=temp.reset_index()['Date'],
open=temp['Open'],
high=temp['High'],
low=temp['Low'],
close=temp['Adj Close'])])

fig.update_layout(title=title,
yaxis_title="USD")

fig.show()
```

Analyze the Stock Market

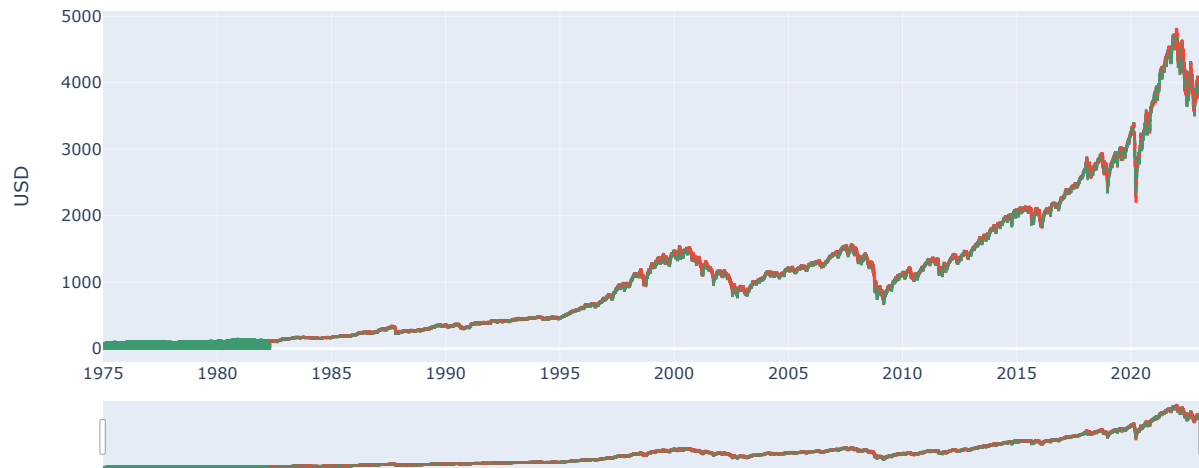
The S&P 500 Index is largely considered an essential benchmark of the stock market. I will analyze it to gain insight about the market

```
In [125]: SP500 = yf.download('^GSPC')

[*****100%*****] 1 of 1 completed
```

```
In [126]: # Generic plot of S&P 500 stock price
candlestick(SP500, 'S&P 500: Market Summary')
```

S&P 500: Market Summary



```
In [127]: def annual_percent_change(df):
# Calculates percent annual return
# Input: df[Date, Open, High, Low, Close, Adj Close]
# Output: df[Year, Open, High, Low, Close, % Change]

years = np.unique(np.array(SP500.index.year))[0:-1]
percent_change = []
year_open, year_close = [], []
year_low, year_high = [], []

for i in years:
    temp = df[(df.index >= str(i)) & (df.index < str(i+1))]['Adj Close']
    percent_change.append(round((temp[-1] / temp[0] - 1) * 100, 2))
    year_open.append(temp[0])
    year_close.append(temp[-1])
    year_low.append(min(temp))
    year_high.append(max(temp))

result = pd.DataFrame({'Date': years,
                        'Open': year_open,
                        'Close': year_close,
                        'Low': year_low,
                        'High': year_high,
                        '% Change': percent_change}).set_index('Date')

return result

SP500_YoY = annual_percent_change(SP500)
```

```
In [128]: average_annual_return = np.mean(SP500_YoY[SP500_YoY.index >= 1950]['% Change'])

SP500_2019 = SP500_YoY[SP500_YoY.index == 2019]['% Change']
SP500_2020 = SP500_YoY[SP500_YoY.index == 2020]['% Change']
SP500_2021 = SP500_YoY[SP500_YoY.index == 2021]['% Change']

SP500_2019, SP500_2020, SP500_2021, average_annual_return
```

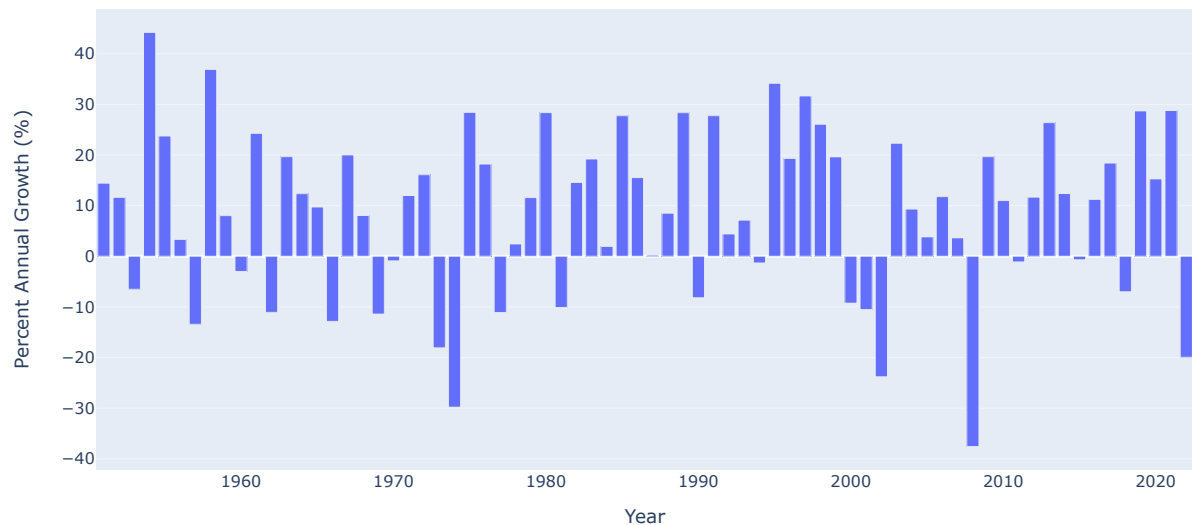
```
Out[128]: (Date
2019      28.71
Name: % Change, dtype: float64,
Date
2020      15.29
Name: % Change, dtype: float64,
Date
2021      28.79
Name: % Change, dtype: float64,
8.894246575342464)
```

The S&P500 index has delivered an average annual growth rate of 8.7% since 1950. However, this rule of thumb did not apply during the pandemic.

- In 2019, S&P 500 grew **29%**
- In 2020, S&P 500 grew **15%**
- In 2021, S&P 500 grew **29%**

```
In [129]: # Plot % Annual Growth
temp = SP500_YoY[SP500_YoY.index > 1950].reset_index() # Only include recent data, (1950 and beyond)
fig = px.bar(temp, x="Date", y="% Change", title='S&P 500: Annual Growth')
fig.update_layout(xaxis_title='Year', yaxis_title='Percent Annual Growth (%)')
```

S&P 500: Annual Growth



```
In [130]: # Calculate percentile
stats.percentileofscore(SP500_YoY['% Change'], SP500_2019), stats.percentileofscore(SP500_YoY['% Change'], SP500_2020), stats.percentileofscore(SP500_YoY['% Change'], SP500_2021)
```

```
Out[130]: (array([90.625]), array([64.58333333]), array([91.66666667]))
```

The plot above shows the percent annual growth in price between 1950 and 2022. We occasionally see long positive stretches, often followed by a big negative year. The 2019-2022 market is certainly bullish, with a growth rate in the **90th, 60th, and 90th percentile** which is reminiscent of the **dot-com bubble in the late 90s**.

II) Modeling and Prediction: one day in the future

- **Objective:** Predict stock prices one day in advance.
- Because future stock prices are very reliant on past prices, I will use the LSTM model which is known for storing past information

Preprocessing

```
In [153]: def split(df, split_ratio):
# Takes in a df and splits it into train and test df

numrows_train = round(split_ratio * df.shape[0])
train_df = df[0:numrows_train]
test_df = df[numrows_train:]
return train_df, test_df
```

```
In [154]: df = SP500
split_ratio = 0.8

train_df, test_df = split(df, split_ratio)
```

```
In [155]: def offset(array, n_features):
# array: an array of normalized training or test data
# n_features: use past n days of data as the feature to predict price
# lookahead: how many days in the future you want to predict
# Output: x_train and y_train
x, y = [], []

for i in range(n_features, len(array)-1):
    x.append(array[i-n_features:i])
    y.append(array[i, 0])

return np.array(x), np.array(y)
```

```
In [156]: scaler = MinMaxScaler(feature_range = (0,1))
n_features = 30 # Use the last 60 days as a feature

# Normalize training and test data
train_scaled = scaler.fit_transform(train_df['Adj Close'].values.reshape(-1,1))
test_scaled = scaler.fit_transform(test_df['Adj Close'].values.reshape(-1,1))

x_train, y_train = offset(train_scaled, n_features)
x_test, y_test = offset(test_scaled, n_features)
```

Build model

Future stock prices are very reliant on past prices, so I will use the LSTM model that can <...>

```
In [157]: def build_model():
K.clear_session()
model = Sequential()

model.add(LSTM(units = 50, return_sequences = True, input_shape = (x_train.shape[1],1)))
model.add(Dropout(0.2))

model.add(LSTM(units = 50, return_sequences = True))
model.add(Dropout(0.2))

model.add(LSTM(units = 50))
model.add(Dropout(0.2))

model.add(Dense(units=1))

return model
```

```
In [158]: model = build_model()
model.compile(loss='mean_squared_error', optimizer='adam')
early_stop = EarlyStopping(monitor='loss', patience=10, verbose=1)
```

Train model

```
In [159]: model = build_model()
model.compile(loss='mean_squared_error', optimizer='adam')

model.fit(x_train, y_train, epochs=6,
        batch_size = 32)
```

```
Epoch 1/6
597/597 [=====] - 10s 13ms/step - loss: 0.0013
Epoch 2/6
597/597 [=====] - 8s 13ms/step - loss: 6.2116e-04
Epoch 3/6
597/597 [=====] - 8s 13ms/step - loss: 5.6431e-04
Epoch 4/6
597/597 [=====] - 7s 13ms/step - loss: 5.2829e-04
Epoch 5/6
597/597 [=====] - 8s 13ms/step - loss: 4.9433e-04
Epoch 6/6
597/597 [=====] - 8s 13ms/step - loss: 4.6595e-04
```

```
Out[159]: <keras.callbacks.History at 0x1b61a8eb250>
```

```
In [160]: ## Runtime Optimization, the model did not improve much past 6 epochs
# epoch, loss = [], []

# for i in range(3, 10): # Optimal num epoch
#     model = build_model()
#     model.compile(loss='mean_squared_error', optimizer='adam')
#     early_stop = EarlyStopping(monitor='loss', patience=10, verbose=1)

#     model.fit(x_train, y_train, epochs=i,
#               batch_size = 32)

#     epoch.append(i)
#     loss.append(model.evaluate(x_test, y_test))
```

Predict

```
In [161]: test_pred = model.predict(x_test)
test_pred = scaler.inverse_transform(test_pred)
```

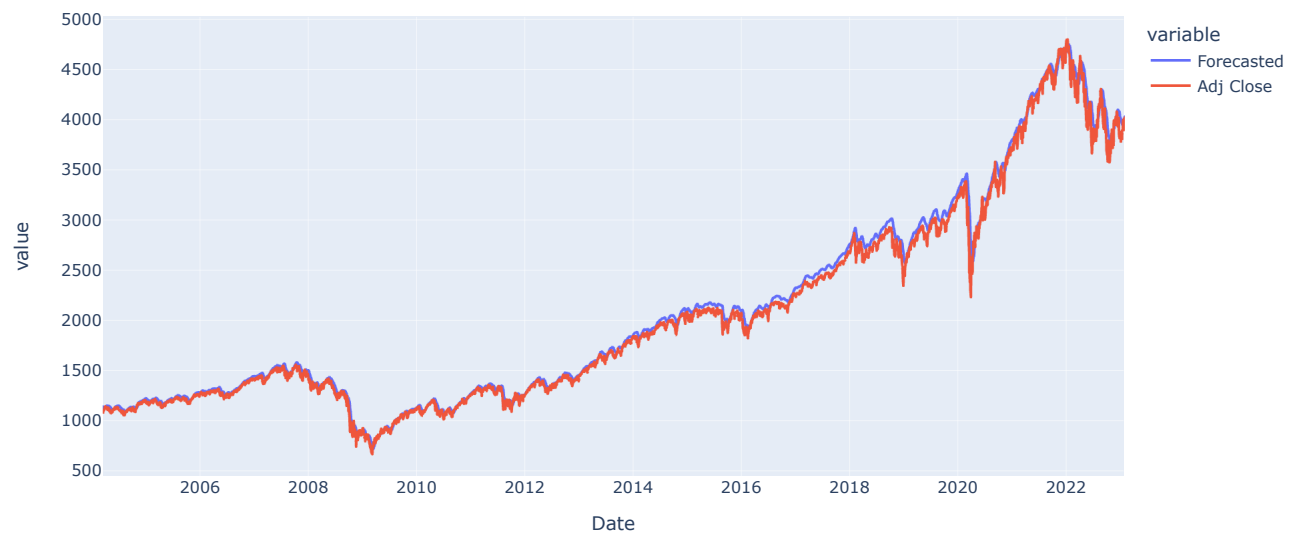
149/149 [=====] - 1s 5ms/step

```
In [162]: summary_df = test_df[n_features:-1].reset_index()
summary_df['Forecasted'] = test_pred
```

Analysis

```
In [163]: fig = px.line(summary_df, x="Date", y=["Forecasted", 'Adj Close'], title='S&P 500: Summary and 1 Day Forecast')
fig.show()
```

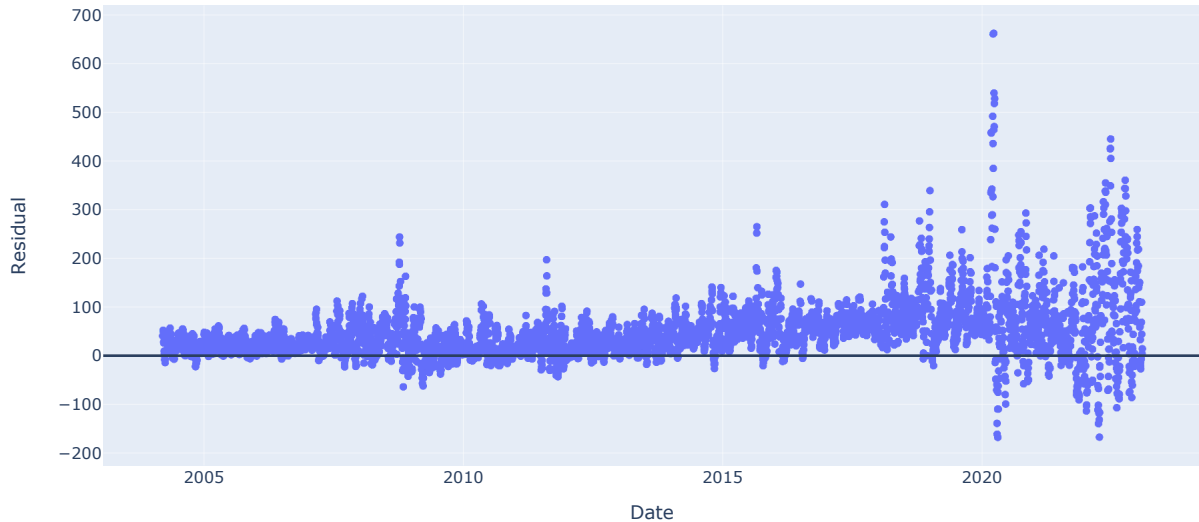
S&P 500: Summary and 1 Day Forecast



The plot tells us that the forecasted price (blue) and true price (red) is overlapped. There are small deviations on a day-to-day basis once you zoom in, but the **model was good at predicting market trends overall.**

```
In [164]: # Residual Plot
summary_df['Residual'] = summary_df['Forecasted'] - summary_df['Adj Close']
fig = px.scatter(summary_df,
                 x='Date', y='Residual', title = 'S&P500: Residual Plot')
fig.add_hline(y=0)
fig.show()
```

S&P500: Residual Plot

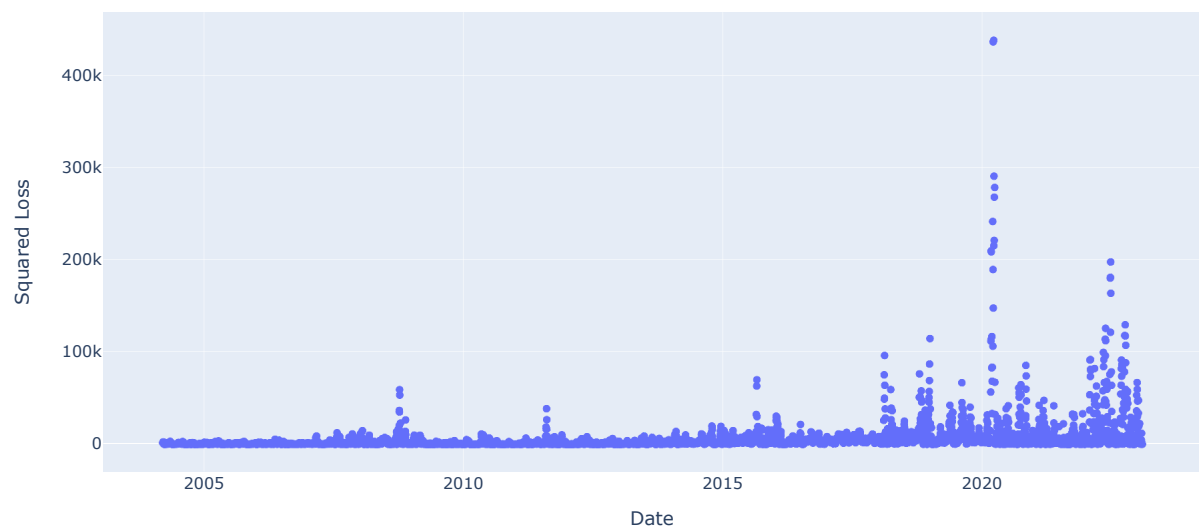


```
In [165]: # RMSE
summary_df['Squared Loss'] = (summary_df['Forecasted'] - summary_df['Adj Close'])**2
RMSE_pre = np.sqrt(summary_df['Squared Loss'][summary_df['Date']<'2018'].mean())
RMSE_post = np.sqrt(summary_df['Squared Loss'][summary_df['Date']>='2018'].mean())
RMSE_pre, RMSE_post
```

Out[165]: (47.29226617199417, 126.4915178773102)

```
In [166]: # Plot Square Loss
px.scatter(summary_df, x='Date', y='Squared Loss', title = 'S&P500: Squared Loss')
```

S&P500: Squared Loss



```
In [167]: # Ten Losses
summary_df.sort_values('Squared Loss', ascending=False)[0:10]
```

```
Out[167]:
```

	Date	Open	High	Low	Close	Adj Close	Volume	Forecasted	Residual	Squared Loss
4024	2020-03-16	2508.590088	2562.979980	2380.939941	2386.129883	2386.129883	7805450000	3048.278564	662.148682	438440.876598
4022	2020-03-12	2630.860107	2660.949951	2478.860107	2480.639893	2480.639893	8850810000	3141.513672	660.873779	436754.152162
4026	2020-03-18	2436.500000	2453.570068	2280.520020	2398.100098	2398.100098	8799300000	2937.523438	539.423340	290977.539568
4029	2020-03-23	2290.709961	2300.729980	2191.860107	2237.399902	2237.399902	7411380000	2765.363525	527.963623	278745.587261
4028	2020-03-20	2431.939941	2453.010010	2295.560059	2304.919922	2304.919922	9053950000	2822.830078	517.910156	268230.929947
4019	2020-03-09	2863.889893	2863.889893	2734.429932	2746.560059	2746.560059	8441290000	3238.299316	491.739258	241807.497674
4027	2020-03-19	2393.479980	2466.969971	2319.780029	2409.389893	2409.389893	7956100000	2879.835938	470.446045	221319.481183
4025	2020-03-17	2425.659912	2553.929932	2367.040039	2529.189941	2529.189941	8370250000	2993.483154	464.293213	215568.187536
4012	2020-02-27	3062.540039	3097.070068	2977.389893	2978.760010	2978.760010	7064710000	3436.890625	458.130615	209883.660615
4013	2020-02-28	2916.899902	2959.719971	2855.840088	2954.219971	2954.219971	8569570000	3411.280029	457.060059	208903.897162

Thoughts

- The above scatterplot plots the **residual** and **squared loss** of our model respectfully. Notice that the model has high loss everytime there is a economic turmoil (like in the late 90s and 2009s). However, the model unequivocally had the **highest loss during lockdown**. Although the model was able to predict trends even during the lockdown, it appears like the model **failed to predict the volatility of the market especially as the beginning of the pandemic**.
- From the table, we can see that the **ten highest model loss occurred between February and March of 2020**.

Other notes

- When we used an 80:20 train-test split, the results were slightly better than 65:35. It might be because the training data learned from the dot-com bubble and was better equipped to predict the pandemic market.
- In general, the residual plots tended to overvalue. However, the spread appears random, which means the error results from noise and not systematic error.

Part 2) Tech Stocks: *Debug

I want to see how the model would perform for the tech sector specifically, seeing how growth in the tech sector was even more pronounced during the pandemic. NASDAQ 100 Technology Sector (NDXT) is an index composed of tech companies like Alphabet, Apple, and Meta (Facebook).

```
In [183]: # df = yf.download('^NDXT')
```

```
In [184]: # split_ratio = 0.75
# train_df, test_df = split(df, split_ratio)
# scaler = MinMaxScaler(feature_range = (0,1))
```

```
In [185]: # n_features = 60

# # Normalize training and test data
# train_scaled = scaler.fit_transform(train_df['Adj Close'].values.reshape(-1,1))
# test_scaled = scaler.fit_transform(test_df['Adj Close'].values.reshape(-1,1))

# x_train, y_train = offset(train_scaled, n_features)
# x_test, y_test = offset(test_scaled, n_features)
```

```
In [186]: # model = build_model()
# model.compile(loss='mean_squared_error', optimizer='adam')

# model.fit(x_train, y_train, epochs=6,
#           batch_size = 32)
```

```
In [187]: # test_pred = model.predict(x_test)
# test_pred = scaler.inverse_transform(test_pred)
```

```
In [188]: # summary_df = test_df[n_features:-1].reset_index()
# summary_df['Forecasted'] = test_pred
```

```
In [189]: # fig = px.line(summary_df, x="Date", y=["Forecasted", 'Adj Close'], title='NDXT: Summary and Forecast')
# fig.show()
```

```
In [190]: # # RMSE
# summary_df['Squared Loss'] = (summary_df['Forecasted'] - summary_df['Adj Close'])**2
# RMSE = np.sqrt(summary_df['Squared Loss'].mean())
# RMSE
```

```
In [182]: # # Plot Error
# summary_df['Residual'] = summary_df['Forecasted'] - summary_df['Adj Close']
# fig = px.scatter(summary_df[summary_df['Date'] > '2005'],
#                 x='Date', y='Residual', title = 'S&P500: Residual Scatterplot')
# fig.add_hline(y=0)
# fig.show()
```

Part 3) Apple (AAPL)

One of the most robust stocks in the tech industry, AAPL stocks tripled in value during the pandemic.

```
In [168]: df = yf.download('AAPL')
```

```
[*****100%*****] 1 of 1 completed
```

```
In [169]: split_ratio = 0.65
train_df, test_df = split(df, split_ratio)
scaler = MinMaxScaler(feature_range = (0,1))
```

```
In [170]: n_features = 60

# Normalize training and test data
train_scaled = scaler.fit_transform(train_df['Adj Close'].values.reshape(-1,1))
test_scaled = scaler.fit_transform(test_df['Adj Close'].values.reshape(-1,1))

x_train, y_train = offset(train_scaled, n_features)
x_test, y_test = offset(test_scaled, n_features)
```

```
In [171]: model = build_model()
model.compile(loss='mean_squared_error', optimizer='adam')

model.fit(x_train, y_train, epochs=6,
          batch_size = 32)
```

```
Epoch 1/6
214/214 [=====] - 8s 23ms/step - loss: 0.0012
Epoch 2/6
214/214 [=====] - 5s 23ms/step - loss: 6.2808e-04
Epoch 3/6
214/214 [=====] - 5s 23ms/step - loss: 6.3165e-04
Epoch 4/6
214/214 [=====] - 5s 23ms/step - loss: 4.4638e-04
Epoch 5/6
214/214 [=====] - 5s 23ms/step - loss: 4.6732e-04
Epoch 6/6
214/214 [=====] - 5s 23ms/step - loss: 5.0887e-04
```

```
Out[171]: <keras.callbacks.History at 0x1b61a9c1c70>
```

```
In [172]: test_pred = model.predict(x_test)
test_pred = scaler.inverse_transform(test_pred)
```

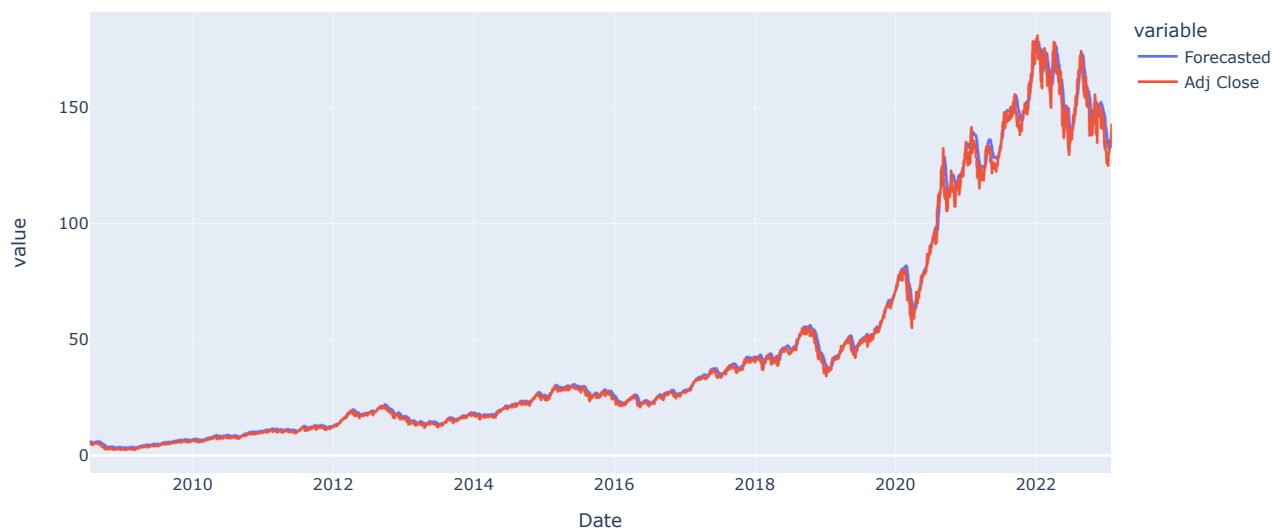
```
115/115 [=====] - 1s 8ms/step
```

```
In [173]: summary_df = test_df[n_features:-1].reset_index()
summary_df['Forecasted'] = test_pred
```



```
In [174]: fig = px.line(summary_df, x="Date", y=["Forecasted", 'Adj Close'], title='AAPL: Summary and Forecast')
fig.show()
```

AAPL: Summary and Forecast

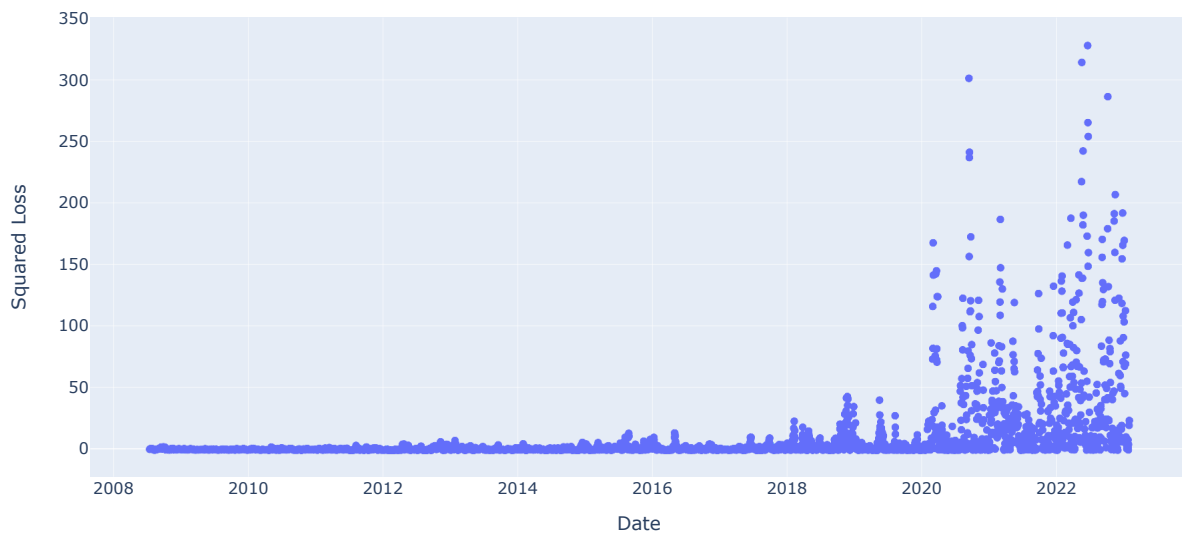


```
In [178]: # RMSE
summary_df['Residual'] = summary_df['Forecasted'] - summary_df['Adj Close']
summary_df['Squared Loss'] = (summary_df['Forecasted'] - summary_df['Adj Close'])**2
RMSE = np.sqrt(summary_df['Squared Loss'].mean())
RMSE
```

```
Out[178]: 2.897211348505621
```

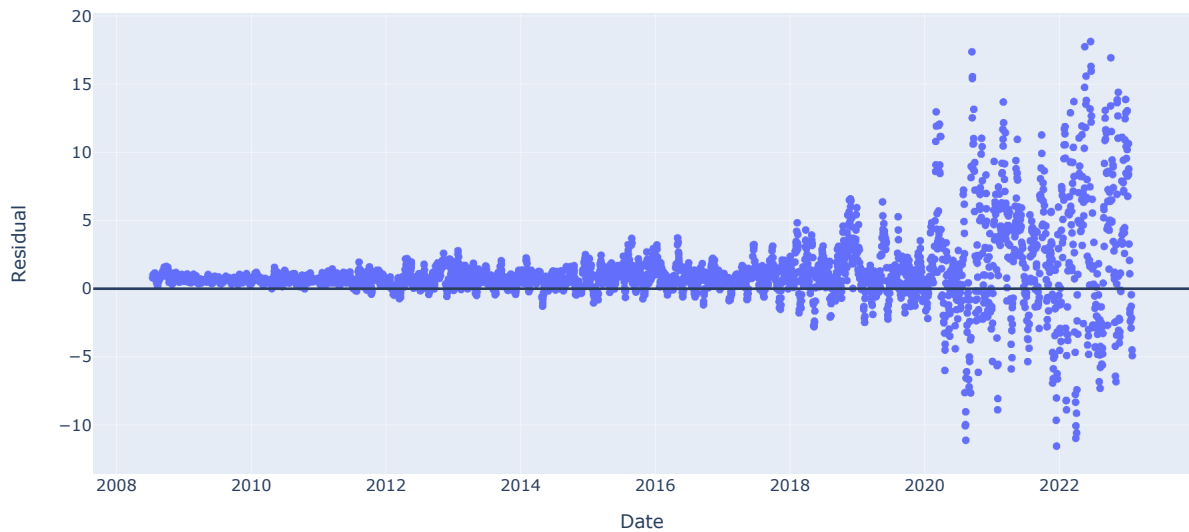
```
In [179]: # Plot Square Error
px.scatter(summary_df, x='Date', y='Squared Loss', title = 'AAPL: Squared Loss')
```

AAPL: Squared Loss



```
In [180]: # Residual Plot
fig = px.scatter(summary_df,
                 x='Date', y='Residual', title = 'AAPL: Residual Plot')
fig.add_hline(y=0)
fig.show()
```

AAPL: Residual Plot



The plots are very similar to S&P500. While the model is good at predicting overall trends, it had difficulty forecasting patterns during the pandemic.

Conclusion

During normal times, the model was a fantastic indicator of the market, with very little loss and predicting the market trends quite accurately. However, during the pandemic, we saw inflated loss values, especially at the onset of the lockdown. Although the model remains a good indicator on average, its predictions should be taken with a grain of salt during irregular times.